

## IX. SUPPLEMENTARY MATERIAL

### A. Traffic distributions

Fig. 15 includes FB\_Hadoop and WebSearch, the two public traffic distributions we use in our experiments. These two distributions are widely used in research [34], [11].




Fig. 15: The traffic workloads used in Escape evaluation.

### B. DPDK Implementation

Listing 1 is the API introduced and implemented to support frame extraction based on a given key, in DPDK `struct rte_ring`. The source code of this implementation is publicly available [14].

```
typedef uint16_t (*comp_obj)
(const void *obj1, const void *obj2);
static __rte_always_inline void *
rte_ring_sc_jump(struct rte_ring *ring,
                 comp_obj func,
                 const void *obj2);
```

Listing 1: DPDK API introduced to support frame extraction.

### C. Micro-Benchmarks

In § VII-A, we use the topology in Fig. 6 to evaluate Escape in terms of its ability to clear deadlocks. Here we have the link status observations from that experiment. Without Escape,  $l_1$ ,  $l_2$ , and  $l_3$  are in XOFF status after  $\sim 18$  ms as a result of deadlock formation, and remain so henceforth (Fig. 16a). Furthermore,  $q_1^I$ ,  $q_2^I$ , and  $q_3^I$  remain within XOFF and XON thresholds (Fig. 8a), indicating PFC activation on corresponding upstream egress interfaces, causing deadlock. In contrast, with Escape enabled,  $l_1$ ,  $l_2$ , and  $l_3$  do not simultaneously remain in XOFF (Fig. 16b).




Fig. 16: Deadlock clearance in Fig. 6 with Escape.

### D. Datacenter Simulations

We evaluate Escape on a scenario mimicking a modern datacenter topology, load, and traffic. We use HPCC [11] for

congestion control. We use the three-level fat-tree topology depicted in Fig. 17a with 8 edge switches, 8 aggregation switches, and 4 core switches. It has 8 source nodes ( $h_1, h_3, h_5, \dots, h_{15}$ ) and 8 destination nodes ( $h_2, h_4, h_6, \dots, h_{16}$ ) connected to the edge switches. Each interconnection is a 40 Gbps link. Each source node starts 10 concurrent data flows. The flows from  $h_1$  and  $h_3$  are destined to  $h_{16}$  and create incast congestion on  $s_6$ . Similarly, the flows from  $h_{13}$  and  $h_{15}$  are destined to  $h_4$ , creating incast congestion on  $s_3$ . The flows from  $h_7$  and  $h_{11}$  are destined to  $h_{12}$  and  $h_8$ , respectively, and create little or no congestion. Similarly, the flows from  $h_5$  and  $h_9$  are destined to  $h_{10}$  and  $h_6$ , respectively, creating little or no congestion. As the second step of this experiment, we disable the links  $s_1 \rightarrow s_6$  and  $s_2 \rightarrow s_3$ , causing the flows destined to  $h_{16}$  and  $h_6$  to take “bouncy” paths. As shown in Fig. 17b the flows from  $h_1$  and  $h_3$  take the new route  $s_3 \rightarrow s_1 \rightarrow s_5 \rightarrow s_2 \rightarrow s_6 \rightarrow h_{16}$  and the flows from  $h_{13}$  and  $h_{15}$  the route  $s_6 \rightarrow s_2 \rightarrow s_4 \rightarrow s_1 \rightarrow s_3 \rightarrow h_3$ . We run the two steps with Escape disabled and then with it enabled. We record link utilization  $l_i$  at each source node  $h_i$  to observe network stability. When Escape is enabled, we additionally record the size of ingress queue  $q_j^I$  of node  $s_j$  on the deadlock-prone path  $s_1 \rightarrow s_5 \rightarrow s_2 \rightarrow s_4 \rightarrow s_1$ .

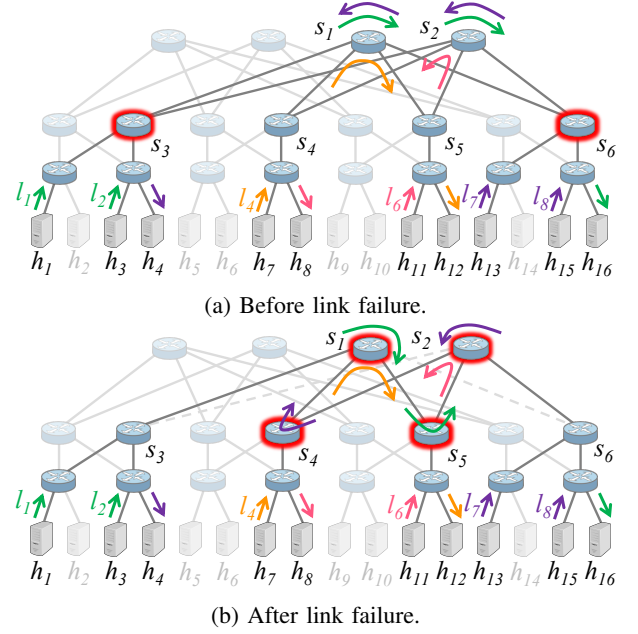


Fig. 17: Three-level fat-tree topology used in the datacenter simulations.

Fig. 18 shows the link utilization before link failure. The flows from  $h_1$  and  $h_3$  go through  $s_6$ , equally sharing the bottleneck link bandwidth, and as a result, the utilization of  $l_1$  (and  $l_2$ ) stabilizes at  $\sim 20$  Gbps. Further, the throughput attained by each of the 10 competing flows on the link, stabilizes at  $\sim 2$  Gbps. Similarly, the flows from  $h_{13}$  and  $h_{15}$  go through the bottleneck link at  $s_3$ , and the utilization of  $l_7$  (and  $l_8$ ) stabilizes at  $\sim 20$  Gbps. The flows from  $h_7$  and  $h_{11}$  do not create congestion, hence  $l_4$  and  $l_6$  are fully utilized.

Without Escape, switches  $s_1$ ,  $s_2$ ,  $s_4$ , and  $s_5$  instantly run into a deadlock in step two (see Fig. 17b), causing the flows going through these switches to halt. With Escape, the link failure does not cause a deadlock and each sender attains its fair share link bandwidth of  $\sim 13.3$  Gbps, as Fig. 19a shows.




Fig. 18: System stability (link utilization) before link failure (HPCC + Escape).

To understand the impact of Escape on system stability under normal conditions and when there is a potential deadlock, we observed ingress queue depths at congestion points and instantaneous bandwidths attained by the completing senders. Fig. 18 shows that throughput convergence of competing senders ( $h_1$ ,  $h_3$  and  $h_{13}$ ,  $h_{15}$ ) is as expected in a network with HPCC. The system converges in  $\sim 5$  ms. As a result of the link failure, the nodes on the deadlock-prone loop become the new congestion hot spots resulting in a different set of competing senders ( $h_1$ ,  $h_3$ ,  $h_7$ , and  $h_{11}$ ,  $h_{13}$ ,  $h_{15}$ ). The system converges and each competing sender attains its fair share in  $\sim 20$  ms (Fig. 19a). At the same time, the queues drain (Fig. 19b) indicating system stability. We can conclude that Escape clears deadlocks caused by HoL blocking in practical fat-tree topologies, while preserving network stability.




Fig. 19: System stability after link failure (HPCC + Escape).